
```
from util import manhattanDistance
from game import Directions
import random, util
```

```
from game import Agent
```

```
class ReflexAgent(Agent):
```

```
    """
```

```
    A reflex agent chooses an action at each choice point by examining
    its alternatives via a state evaluation function.
```

```
    The code below is provided as a guide. You are welcome to change
    it in any way you see fit, so long as you don't touch our method
    headers.
```

```
    """
```

```
    def __init__(self):
```

```
        self.lastPositions = []
```

```
        self.dc = None
```

```
    def getAction(self, gameState):
```

```
        """
```

```
        getAction chooses among the best options according to the evaluation
        function.
```

```
        getAction takes a GameState and returns some Directions.X for some X in the
        set {North, South, West, East, Stop}
```

```
-----
    Description of GameState and helper functions:
```

```
    A GameState specifies the full game state, including the food, capsules,
    agent configurations and score changes. In this function, the |gameState|
    argument
```

```
    is an object of GameState class. Following are a few of the helper methods
    that you
```

```
    can use to query a GameState object to gather information about the present
    state
    of Pac-Man, the ghosts and the maze.
```

```
    gameState.getLegalActions():
```

```
        Returns the legal actions for the agent specified. Returns Pac-Man's
    legal moves by default.
```

```
    gameState.generateSuccessor(agentIndex, action):
```

```
        Returns the successor state after the specified agent takes the action.
        Pac-Man is always agent 0.
```

```
    gameState.getPacmanState():
```

```
        Returns an AgentState object for pacman (in game.py)
        state.configuration.pos gives the current position
        state.direction gives the travel vector
```

```
    gameState.getGhostStates():
```

```
        Returns list of AgentState objects for the ghosts
```

```
    gameState.getNumAgents():
```

```
        Returns the total number of agents in the game
```

```
    gameState.getScore():
```

```
        Returns the score corresponding to the current state of the game
```

The GameState class is defined in pacman.py and you might want to look into

```

that for
    other helper methods, though you don't need to.
    """
    # Collect legal moves and successor states
    legalMoves = gameState.getLegalActions()

    # Choose one of the best actions
    scores = [self.evaluationFunction(gameState, action) for action in
legalMoves]
    bestScore = max(scores)
    bestIndices = [index for index in range(len(scores)) if scores[index] ==
bestScore]
    chosenIndex = random.choice(bestIndices) # Pick randomly among the best

    return legalMoves[chosenIndex]

def evaluationFunction(self, currentGameState, action):
    """
    The evaluation function takes in the current and proposed successor
    GameStates (pacman.py) and returns a number, where higher numbers are
    better.

    The code below extracts some useful information from the state, like the
    remaining food (oldFood) and Pacman position after moving (newPos).
    newScaredTimes holds the number of moves that each ghost will remain
    scared because of Pacman having eaten a power pellet.
    """
    # Useful information you can extract from a GameState (pacman.py)
    successorGameState = currentGameState.generatePacmanSuccessor(action)
    newPos = successorGameState.getPacmanPosition()
    oldFood = currentGameState.getFood()
    newGhostStates = successorGameState.getGhostStates()
    newScaredTimes = [ghostState.scaredTimer for ghostState in newGhostStates]

    return successorGameState.getScore()

def scoreEvaluationFunction(currentGameState):
    """
    This default evaluation function just returns the score of the state.
    The score is the same one displayed in the Pacman GUI.

    This evaluation function is meant for use with adversarial search agents
    (not reflex agents).
    """
    return currentGameState.getScore()

class MultiAgentSearchAgent(Agent):
    """
    This class provides some common elements to all of your
    multi-agent searchers. Any methods defined here will be available
    to the MinimaxPacmanAgent, AlphaBetaPacmanAgent & ExpectimaxPacmanAgent.

    You *do not* need to make any changes here, but you can if you want to
    add functionality to all your adversarial search agents. Please do not
    remove anything, however.

    Note: this is an abstract class: one that should not be instantiated. It's
    only partially specified, and designed to be extended. Agent (game.py)
    is another abstract class.
    """

    def __init__(self, evalFn = 'scoreEvaluationFunction', depth = '2'):

```

```

self.index = 0 # Pacman is always agent index 0
self.evaluationFunction = util.lookup(evalFn, globals())
self.depth = int(depth)

#####
#####
# Problem 1b: implementing minimax

class MinimaxAgent(MultiAgentSearchAgent):
    """
    Your minimax agent (problem 1)
    """

    def getAction(self, gameState):
        """
        Returns the minimax action from the current gameState using self.depth
        and self.evaluationFunction. Terminal states can be found by one of the
        following:
        pacman won, pacman lost or there are no legal moves.

        Here are some method calls that might be useful when implementing minimax.

        gameState.getLegalActions(agentIndex):
            Returns a list of legal actions for an agent
            agentIndex=0 means Pacman, ghosts are >= 1

        gameState.generateSuccessor(agentIndex, action):
            Returns the successor game state after an agent takes an action

        gameState.getNumAgents():
            Returns the total number of agents in the game

        gameState.getScore():
            Returns the score corresponding to the current state of the game

        gameState.isWin():
            Returns True if it's a winning state

        gameState.isLose():
            Returns True if it's a losing state

        self.depth:
            The depth to which search should continue
        """

        # BEGIN_YOUR_CODE (our solution is 20 lines of code, but don't worry if you
        deviate from this)
        def getMin(gameState, index, d):
            valAction = (float("inf"), Directions.STOP)
            nextAgent = index + 1
            depth = d
            if nextAgent == gameState.getNumAgents():
                nextAgent = 0
                depth -= 1
            for action in gameState.getLegalActions(index):
                minimaxVal = Vminimax(gameState.generateSuccessor(index, action),
                nextAgent, depth)
                if minimaxVal < valAction[0]: # update expected utility
                    valAction = (minimaxVal, action)
            return valAction

        def getMax(gameState, index, d):
            valAction = (float("-inf"), Directions.STOP)

```

```

        for action in gameState.getLegalActions(index):
            minimaxVal = Vminimax(gameState.generateSuccessor(index, action), index
+ 1, d)
            if minimaxVal > valAction[0]: # update expected utility
                valAction = (minimaxVal, action)
            return valAction

def Vminimax(gameState, index, d):
    if d == 0 or gameState.isLose() or gameState.isWin():
        return self.evaluationFunction(gameState)
    if index == 0: # pacman
        return getMax(gameState, index, d)[0]
    else: # ghost
        return getMin(gameState, index, d)[0]

return getMax(gameState, self.index, self.depth)[1]
# END_YOUR_CODE

#####
#####
# Problem 2a: implementing alpha-beta

class AlphaBetaAgent(MultiAgentSearchAgent):
    """
    Your minimax agent with alpha-beta pruning (problem 2)
    """

    def getAction(self, gameState):
        """
        Returns the minimax action using self.depth and self.evaluationFunction
        """

        # BEGIN_YOUR_CODE (our solution is 36 lines of code, but don't worry if you
deviate from this)
        def getMin(gameState, agentIndex, d, alpha, beta):
            valAction = (float("inf"), Directions.STOP)
            nextAgent = agentIndex + 1
            depth = d
            if nextAgent == gameState.getNumAgents():
                nextAgent = 0
                depth -= 1
            for action in gameState.getLegalActions(agentIndex):
                minimaxVal = Vminimax(gameState.generateSuccessor(agentIndex, action),
nextAgent, depth, alpha, beta)
                if minimaxVal < valAction[0]: # update expected utility
                    valAction = (minimaxVal, action)
                beta = min(beta, valAction[0])
                if valAction[0] < alpha:
                    return valAction
            return valAction

        def getMax(gameState, agentIndex, d, alpha, beta):
            valAction = (float("-inf"), Directions.STOP)
            for action in gameState.getLegalActions(agentIndex):
                minimaxVal = Vminimax(gameState.generateSuccessor(agentIndex, action),
agentIndex + 1, d, alpha, beta)
                if minimaxVal > valAction[0]: # update expected utility
                    valAction = (minimaxVal, action)
                alpha = max(alpha, valAction[0])
                if valAction[0] > beta:
                    return valAction
            return valAction

        def Vminimax(gameState, agentIndex, d, alpha, beta):

```

```

        if d == 0 or gameState.isLose() or gameState.isWin():
            return self.evaluationFunction(gameState)
        if agentIndex == 0: # pacman
            return getMax(gameState, agentIndex, d, alpha, beta)[0]
        else: # ghost
            return getMin(gameState, agentIndex, d, alpha, beta)[0]

    return getMax(gameState, self.index, self.depth, float("-inf"),
float("inf"))[1]
# END_YOUR_CODE

#####
#####
# Problem 3b: implementing expectimax

class ExpectimaxAgent(MultiAgentSearchAgent):
    """
    Your expectimax agent (problem 3)
    """

    def getAction(self, gameState):
        """
        Returns the expectimax action using self.depth and self.evaluationFunction

        All ghosts should be modeled as choosing uniformly at random from their
        legal moves.
        """

        # BEGIN_YOUR_CODE (our solution is 20 lines of code, but don't worry if you
        deviate from this)
        def getExpectimax(gameState, agentIndex, d):
            valAction = []
            nextAgent = agentIndex + 1
            depth = d
            if nextAgent == gameState.getNumAgents():
                nextAgent = 0
                depth -= 1
            for action in gameState.getLegalActions(agentIndex):
                minimaxVal = Vexpectimax(gameState.generateSuccessor(agentIndex,
action), nextAgent, depth)
                valAction.append(minimaxVal)
            return sum(valAction)/len(valAction)

        def getMax(gameState, agentIndex, d):
            valAction = (float("-inf"), Directions.STOP)
            for action in gameState.getLegalActions(agentIndex):
                minimaxVal = Vexpectimax(gameState.generateSuccessor(agentIndex,
action), agentIndex + 1, d)
                if minimaxVal > valAction[0]:
                    valAction = (minimaxVal, action)
            return valAction

        def Vexpectimax(gameState, agentIndex, d):
            if d == 0 or gameState.isLose() or gameState.isWin():
                return self.evaluationFunction(gameState)
            if agentIndex == 0: # pacman
                return getMax(gameState, agentIndex, d)[0]
            else: # ghost
                return getExpectimax(gameState, agentIndex, d)

        return getMax(gameState, self.index, self.depth)[1]
# END_YOUR_CODE

#####

```

```

#####
# Problem 4a (extra credit): creating a better evaluation function

def betterEvaluationFunction(currentGameState):
    """
    Your extreme, unstoppable evaluation function (problem 4). Note that you
    can't fix a seed in this function.

    DESCRIPTION: <write something here so we know what you did>
    """

    # BEGIN_YOUR_CODE (our solution is 13 lines of code, but don't worry if you
    deviate from this)
    curPos = currentGameState.getPacmanPosition()
    foodList = currentGameState.getFood().asList()
    ghostStates = currentGameState.getGhostStates()
    capsules = currentGameState.getCapsules()
    score = currentGameState.getScore()
    minGhostDistance = float("-inf")
    minCapsuleDistance = float("-inf")
    scaredScore = 0

    # Iterate over ghost states
    distanceScores = []
    ghostDistances = []
    for ghost in ghostStates:
        dist = util.manhattanDistance(ghost.getPosition(), curPos)
        ghostDistances.append(dist)
        if ghost.scaredTimer == 0:
            distanceScores.append((dist + 10) * -1)
            break
        if ghost.scaredTimer > 8:
            if dist <= 4: distanceScores.append(40)
            if dist <= 3: distanceScores.append(60)
            if dist <= 2: distanceScores.append(80)
            if dist <= 1: distanceScores.append(90)
    if len(ghostDistances) > 0: minGhostDistance = min(ghostDistances)
    if len(distanceScores) > 0: scaredScore = max(distanceScores)

    # Iterate over capsules
    capsuleDistances = []
    for caps in capsules: capsuleDistances.append(util.manhattanDistance(caps,
curPos))
    if len(capsuleDistances) > 0: minCapsuleDistance = min(capsuleDistances)

    # Iterate over food list
    foodDistances = []
    for food in foodList: foodDistances.append(util.manhattanDistance(food,
curPos))

    # Update score
    if minCapsuleDistance < minGhostDistance: score += 30
    score += scaredScore
    score -= .35 * sum(foodDistances)

    return score
    # END_YOUR_CODE

# Abbreviation
better = betterEvaluationFunction

```